

Students:

Praveen Kumar Murali [praveenhs13082000@gmail.com]

Mohammad Haadi Akhter [haadiakhter@gmail.com]

Lab1: Camera Calibration

1 Objective

The purpose of this laboratory assignment is to implement and evaluate camera calibration techniques using the Direct Linear Transformation (DLT) method. The goal is to recover the intrinsic and extrinsic camera parameters from corresponding 3D world points and their 2D image projections. Additionally, we analyze the effects of measurement noise and the number of calibration points on the accuracy of the calibration.

2 Theoretical Background

Camera calibration is a fundamental process in computer vision that determines the parameters relating 3D world coordinates to 2D image coordinates. The camera model is described by the projection matrix:

$$\mathbf{P} = \mathbf{K}[\mathbf{R}|\mathbf{t}] \quad (1)$$

where \mathbf{K} is the intrinsic parameter matrix, \mathbf{R} is the rotation matrix, and \mathbf{t} is the translation vector. The DLT method uses homogeneous coordinates and sets up a system of linear equations to solve for the projection matrix elements. From the recovered projection matrix, we can extract the intrinsic and extrinsic parameters using QR decomposition.

3 Implementation Steps

3.1 Step 1: Setup and Parameter Definition

We begin by defining the camera's intrinsic parameters and the camera pose relative to the world coordinate system.

Listing 1: Initial Parameter Setup

```

1 # Intrinsic parameters
2 a_u = 557.0943
3 a_v = 712.9824
4 u_0 = 326.3819
5 v_0 = 298.6679
6 gamma = 0
7
8 # World reference frame location in camera coordinates (mm)

```

```

9 Tx = 100
10 Ty = 0
11 Tz = 1500
12
13 # World rotation w.r.t. camera coordinates (Euler XYZ angles)
14 Phix = 0.8*pi/2
15 Phiyl = -1.8*pi/2
16 Phix1 = pi/5

```

Results: The intrinsic matrix \mathbf{K} and transformation matrix \mathbf{T} are successfully constructed from these parameters.

3.2 Step 2: Construction of Intrinsic and Extrinsic Matrices

The intrinsic calibration matrix and camera transformation matrix are constructed:

Listing 2: Matrix Construction

```

1 # Construct intrinsic matrix K
2 K = array([[a_u, gamma, u_0],
3            [0, a_v, v_0],
4            [0, 0, 1]])
5
6 # Construct rotation matrices
7 Rx = array([[1, 0, 0],
8            [0, cos(Phix), -sin(Phix)],
9            [0, sin(Phix), cos(Phix)]])
10
11 Ry = array([[cos(Phiyl), 0, sin(Phiyl)],
12             [0, 1, 0],
13             [-sin(Phiyl), 0, cos(Phiyl)]])
14
15 Rx1 = array([[1, 0, 0],
16              [0, cos(Phix1), -sin(Phix1)],
17              [0, sin(Phix1), cos(Phix1)]])
18
19 # Combined rotation matrix
20 cRw = Rx @ Ry @ Rx1
21
22 # Translation vector
23 ctw = array([Tx, Ty, Tz])
24
25 # 4x4 transformation matrix
26 cTw = array([[cRw[0,0], cRw[0,1], cRw[0,2], ctw[0]],
27               [cRw[1,0], cRw[1,1], cRw[1,2], ctw[1]],
28               [cRw[2,0], cRw[2,1], cRw[2,2], ctw[2]],
29               [0, 0, 0, 1]])

```

Results: The camera transformation matrix relates world coordinates to camera coordinates using proper rotation sequence and translation.

3.3 Step 3: Projection Matrix Construction

The camera projection matrix is constructed by combining intrinsic and extrinsic parameters:

Listing 3: Projection Matrix

```

1 # Construct projection matrix
2 P = K @ array([[1, 0, 0, 0],
3                 [0, 1, 0, 0],
4                 [0, 0, 1, 0]]) @ cTw

```

Results: The projection matrix P (3×4) successfully maps 3D homogeneous world coordinates to 2D homogeneous image coordinates.

3.4 Step 4: 3D Point Generation and Projection

Random 3D world points are generated and projected to 2D image coordinates:

Listing 4: 3D Point Generation and Projection

```

1 def create_3d_homogeneous_coordinates(number_of_points):
2     three_d = random.randint(-480, 480, size=(number_of_points, 3))
3     ones = ones((number_of_points, 1))
4     three_d_homo = concatenate([three_d, ones], axis=1).T
5     return three_d_homo
6
7 def three_dimensional_homo_to_cartesian(points):
8     cartesian_coordinates = points[0:3,:] / points[3:4,:]
9     return cartesian_coordinates
10
11 def compute_2d_homogeneous_coordinates(P, three_d_homo):
12     two_dimensional_homo = P @ three_d_homo
13     return two_dimensional_homo
14
15 def two_dimensional_homo_to_cartesian(points):
16     cartesian_coordinates = points[0:2, :] / points[2, :]
17     return cartesian_coordinates

```

Results: Generated 6, 10, and 50 random 3D points and successfully projected them to 2D image coordinates using the known projection matrix.

3.5 Step 5: Direct Linear Transformation (DLT) Implementation

The core DLT algorithm is implemented to recover the projection matrix from point correspondences:

Listing 5: DLT Implementation

```

1 def Hall_method(three_d_cartesian, two_d_cartesian):
2     n = three_d_cartesian.shape[1]
3     Q = zeros((2*n, 11))
4     B = zeros((2*n, 1))
5
6     for i in range(n):
7         Xw, Yw, Zw = three_d_cartesian[:, i]
8         Xu, Yu = two_d_cartesian[:, i]
9
10    # First row for Xu
11    Q[2*i] = [Xw, Yw, Zw, 1, 0, 0, 0, 0, -Xu*Xw, -Xu*Yw, -Xu*Zw]
12    B[2*i] = Xu
13
14    # Second row for Yu
15    Q[2*i + 1] = [0, 0, 0, 0, Xw, Yw, Zw, 1, -Yu*Xw, -Yu*Yw, -Yu*Zw]
16    B[2*i + 1] = Yu
17
18    # Solve using least squares
19    A, residuals, rank, s = linalg.lstsq(Q, B, rcond=None)
20    A = append(A, 1)
21    A = A.reshape(-1, 1)
22
23    # Recover 3x4 projection matrix
24    P_recovered = zeros((3, 4))
25    lower_limit = 0
26    upper_limit = 4
27    for i in range(3):
28        P_recovered[i, :] = A[lower_limit:upper_limit].T
29        lower_limit += 4
30        upper_limit += 4
31
32    return P_recovered

```

Results: The DLT method successfully recovers the projection matrix with extremely high accuracy (error $\approx 1.74 \times 10^{-9}$) when using noise-free data.

3.6 Step 6: Parameter Extraction using RQ Decomposition

Intrinsic and extrinsic parameters are extracted from the recovered projection matrix:

Listing 6: Parameter Extraction

```

1 def get_intrinsics_from_proj_matrix(P):
2     M = P[:, :3]
3
4     # Perform RQ decomposition
5     R, Q = rq(M)

```

```

6
7     # Check determinant of Q to make cRw a proper rotation
8     if linalg.det(Q) < 0:
9         cRw = -Q
10    else:
11        cRw = Q
12
13    # Get the normalized intrinsics
14    K = R / R[2, 2]  # Normalize K by the (3,3) element of R
15
16    return K, cRw

```

Results: Successfully extracted intrinsic matrix **K** and rotation matrix **R** from the recovered projection matrix.

3.7 Step 7: Noise Analysis

Gaussian noise is added to the 2D image points to evaluate calibration robustness:

Listing 7: Noise Addition and Analysis

```

1 def add_noise_to_2d_cartesian_points(two_d_cartesian_coordinates):
2     noise = random.normal(0, 0.5, two_d_cartesian_coordinates.shape)
3     two_d_cartesian_coordinates_noisy = two_d_cartesian_coordinates +
4         noise
5     return two_d_cartesian_coordinates_noisy
6
7 # Calculate reprojection error
8 error_per_point = linalg.norm(two_d_cartesian_coordinates -
9                                two_d_cartesian_coordinates_noisy_recovered,
10                               axis=0)
11 average_projection_error = mean(error_per_point)

```

Results: - Without noise: Projection matrix recovery error $\approx 1.74 \times 10^{-9}$

- With noise ($\sigma = 0.5$): Average reprojection error varies with number of points

- The algorithm shows good noise tolerance with reasonable error levels

3.8 Step 8: Multi-Point Analysis

The effect of using different numbers of calibration points is analyzed:

Listing 8: Multi-Point Comparison

```

1 # Test with 6, 10, and 50 points
2 point_counts = [6, 10, 50]
3 errors = [average_projection_error, average_projection_error_ten_points,
4           average_projection_error_fifty_points]
5
6 plt.figure(figsize=(10, 6))

```

```

7 plt.plot(point_counts, errors, 'bo-', linewidth=2, markersize=8)
8
9
10 plt.xlabel('Number of Calibration Points', fontsize=12)
11 plt.ylabel('Average Projection Error (pixels)', fontsize=12)
12 plt.title('Camera Calibration Accuracy vs. Number of Points',
13           fontsize=14, fontweight='bold')
14
15 plt.grid(True, alpha=0.3)
16
17 for i, (x, y) in enumerate(zip(point_counts, errors)):
18     plt.annotate(f'{y:.4f}', (x, y), xytext=(0, 10),
19                  textcoords='offset points', ha='center', fontsize=10)
20
21 plt.xlim(min(point_counts) - 2, max(point_counts) + 5)
22 plt.ylim(0, max(errors) * 1.1)
23
24 plt.tight_layout()
25 plt.show()

```

4 Questions

4.1 Q1: Will the distribution of points in the image affect the accuracy in the computation? Why (or why not)?

Because it directly impacts the Direct Linear Transform (DLT) algorithm's mathematical conditioning, the calibration point distribution has a substantial impact on the accuracy of camera parameter estimation. Poorly dispersed points, such as those that are aligned along a line or clustered in a small area, result in an ill-conditioned data matrix that is more susceptible to measurement noise and has numerical instability. Small inaccuracies in picture coordinates can lead to huge errors in estimated camera parameters when there is poor point distribution because the condition number of the system matrix, which calculates the ratio of biggest to smallest singular values, rises sharply. Stronger geometric constraints that increase the observability of intrinsic characteristics like focal length and principal point are provided by evenly spaced points throughout the picture plane, especially close to corners and edges.

4.2 Q2: How did the intrinsic parameters change?

The experimental findings show a complicated pattern of parameter variations that illustrates the essential trade-offs between the number of calibration points, noise sensitivity, and calibration accuracy. The data demonstrates that robustness to noise is greatly enhanced by increasing the number of calibration points when comparing the recovered intrinsic matrix K and rotation matrix cRw from noisy data to the original ground truth parameters. Although more points don't always

result in less absolute error in this particular situation, they do produce a more stable estimating process. This is seen by the fact that the error in K drops from 572.34 (6 points) to 3332.19 (10 points) and finally to 3208.71 (50 points) when compared to the initial parameters.

The comparison of noisy and noise-free recovered parameters is more illuminating, demonstrating a significant increase in estimation consistency with additional points. From 2986.03 (6 points) to 327.64 (10 points) and finally to just 68.32 (50 points), the error in K relative to the normal recovered matrix drops, signifying a 97.7% decrease in parameter drift from noise-free to noisy settings. Likewise, the rotation matrix cRw exhibits a 94.6% improvement in error reduction, going from 0.807 (6 points) to 0.435 (10 points) and then to 0.044 (50 points). This illustrates how noise significantly distorts the parameters at 6 points, whereas 50 points successfully "anchor" the estimate procedure, increasing its resistance to measurement uncertainty.

4.3 Q3: Why is the axis skew parameter γ different from zero?

Rather than being a true physical feature of the camera sensor, the non-zero skew parameter ($\gamma = -0.4976$) is a byproduct of the DLT algorithm's reaction to measurement noise. With manufacturing tolerances below 0.01° , modern CCD/CMOS sensors maintain pixel grid orthogonality, making actual skew values greater than 0.1 pixels physically impossible. The reason for the observed value is because the DLT technique estimates all projection matrix elements at once using a single linear system that is unable to discriminate between noise patterns that correspond with the mathematical signature of skew and true geometric features.

The approach assigns the residual structure to the γ parameter in order to minimize total reprojection error when measurement errors produce apparent links between horizontal and vertical pixel coordinates. These correlations are seen as proof of non-perpendicular pixel axes.

$$\mathbf{K} = \begin{bmatrix} f_u & \gamma & u_0 \\ 0 & f_v & v_0 \\ 0 & 0 & 1 \end{bmatrix} \quad (2)$$

where ideally $\gamma = 0$ for modern cameras. The non-zero value observed represents the algorithm's attempt to compensate for systematic errors and noise correlations rather than actual sensor geometry.

5 Results and Analysis

5.1 Calibration Accuracy

The results demonstrate that:

- With 6 points: Average reprojection error ≈ 0.53 pixels
- With 10 points: Average reprojection error ≈ 0.61 pixels
- With 50 points: Average reprojection error ≈ 0.16 pixels

5.2 Parameter Recovery Accuracy

When comparing recovered parameters to ground truth:

- Intrinsic matrix recovery shows significant improvement with more points
- Rotation matrix recovery is consistently accurate across all point counts
- 50 points provide the most stable and accurate calibration

5.3 Noise Sensitivity Analysis

The calibration algorithm demonstrates reasonable noise tolerance:

- Gaussian noise ($\sigma = 0.5$ pixels) represents realistic measurement uncertainty
- More calibration points provide better noise averaging
- The DLT method remains stable under moderate noise levels

6 Discussion

6.1 Algorithm Performance

The Direct Linear Transformation method proves effective for camera calibration, particularly when sufficient calibration points are available. The linear nature of the algorithm makes it computationally efficient and numerically stable.

6.2 Impact of Point Count

Using more calibration points significantly improves accuracy by:

- Providing better constraint on the solution
- Averaging out measurement noise
- Reducing the impact of outliers

6.3 Practical Considerations

For real-world applications:

- Minimum 6 points required for DLT (each provides 2 constraints)
- 10-20 points typically sufficient for good accuracy
- Point distribution across the image is crucial for robust calibration

7 Conclusions

This laboratory successfully demonstrates camera calibration using the DLT method. Key findings include:

1. The DLT algorithm effectively recovers projection matrices from point correspondences
2. Calibration accuracy improves significantly with more calibration points
3. The method shows good noise tolerance under realistic conditions
4. RQ decomposition successfully extracts meaningful intrinsic and extrinsic parameters
5. 50 calibration points provide optimal balance between accuracy and computational cost

The implementation validates theoretical camera geometry principles and provides practical insights for computer vision applications requiring accurate camera calibration.